

# Project 2: TD methods and n-step bootstrapping

August Jonasson

February 25, 2025

## Part 1: Expected SARSA and importance sampling ratio

Suppose environment consists of four states  $\{s_0, s_1, s_2, s_3\}$ , with an action space  $\{+, -\}$  and deterministic dynamics  $p(s_{i+1}|s_i, +) = p(s_{i-1}|s_i, -) = 1$ , with exceptions  $p(s_0|s_3, +) = p(s_3|s_0, -) = 1$ .

### Task 1

Given policy  $\pi(+) = \pi(-) = 0.5$ ,  $\forall s$ , and having observed the transition

$$(s_0, +) \rightarrow (s_1, r = 2),$$

we want to compute the one-step *expected SARSA* update of the action-value estimate  $Q(s_0, +)$ . In order to do this, we assume initial estimates  $Q(s, a) = 1$ ,  $\forall s, a$ , learning rate  $\alpha = 0.2$ , and discount factor  $\gamma = 0.9$ . Now, using equation (6.9) from the course book (Sutton 2018, p. 133)

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

with the given policy, parameters, values and observations, we get the one-step update

$$Q(s_0, +) \leftarrow 1 + 0.2 [2 + 0.9(0.5 + 0.5) - 1] = 1.38,$$

and we are done.

### Task 2

Suppose an agent is following a behavior policy  $b(+) = 0.7$  and  $b(-) = 0.3$ , with the target policy equal to that of the previous task. Having observed the trajectory

$$(s_0, +) \rightarrow (s_1, +) \rightarrow s_3,$$

we want to compute the *importance sampling ratio*  $\rho_{0:1}$ . In order to do this, we simply use equation (5.3) from the course book (p. 104), which defines

$$\rho_{t:T-1} := \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}.$$

Plugging in our policies for the two time steps, we get

$$\rho_{0:1} = \left(\frac{0.5}{0.7}\right)^2 \approx 0.51,$$

which can be interpreted as: “the observed trajectory is roughly half as likely to occur under the target policy than under the behavior policy”, and we are done.

## Part 2: TD methods applied to gridworld with a monster

In this task, an agent acts in an  $N^2$  gridworld with a randomly moving monster and a randomly spawned, stationary apple. The goal of the agent is to collect the apple as many times as possible before the episode is over, without getting caught by the monster.

Each state is represented by three tuples (agent position, monster position and apple position) and the entire space is given by

$$\mathcal{S} = \{(x_p, y_p), (x_m, y_m), (x_a, y_a) \mid x_p, y_p, x_m, y_m, x_a, y_a \in \{0, \dots, N-1\}\},$$

i.e.  $|\mathcal{S}| = N^6$ . The action space is

$$\mathcal{A} = \{\text{left, up, right, down}\},$$

which yields that the number of state-action pairs is  $4 \cdot N^6$ .

Some dynamics of the system are:

- the agent moves deterministically according to chosen action; if a wall is hit, the agent instead remains in place,
- the monster moves uniformly random in all four directions; a wall hit is the same as for the agent,
- the apple, if collected by agent, randomly respawns in a new empty cell,
- an episode ends if the monster catches the agent, or after  $T$  time steps.

Given a state-action pair, the rewards are

$$R(s, a) = \begin{cases} +1 & \text{if the agent collects an apple,} \\ -1 & \text{if the agent is caught by the monster,} \\ 0 & \text{otherwise.} \end{cases}$$

For the following tasks, we use hyperparameters:

- side length  $N = 5$ ,
- episode length  $T = 30$ ,
- discount factor  $\gamma = 0.9$ ,
- learning rate  $\alpha = 0.1$  and
- all state-action value estimates  $Q(s, a)$  initialized at zero.

## Task 1: write code

In this first task we implement the gridworld system as defined above, as well as the following TD methods, using an  $\epsilon$ -greedy policy, for action-value estimation: SARSA, off-policy Q-learning, Double Q-learning, and  $n$ -step SARSA.

See Appendix 1 for the Python code.

## Task 2: plot the learning curves

In this task we are to answer which method converges the fastest and we also want to decide which  $n$  is optimal, in the  $n$ -step SARSA method.




Figure 1: Moving average reward per episode (each average computed with previous 1000 episodes) for the different TD methods using  $\epsilon$ -greedy policy with exponentially decaying  $\epsilon$  as a function of the current number of episodes (such that  $\epsilon_1 = 1.0$  and  $\epsilon_{10^5} \approx 0.01$ ).

From Figure 1 we see that all methods have similar convergence rates. For all five of them, convergence seems to happen after somewhere between 40 and 60 thousand episodes.

Considering the  $n$ -step SARSA methods, as we already mentioned, they converge at similar rates. Thus, if we are to answer which one is optimal we look to the one that produces the highest average reward, which is  $n = 1$ , and we are done.

### Task 3: discussion

In this last task, we discuss the following:

1. why all three  $n$ -step SARSA methods have similar rates of convergence,
2. why the apparent differences between regular Q-learning and Double Q-learning,
3. how convergence would be affected if the monster instead moves deterministically towards the agent.

1)

There is a case to be made for the bias-variance trade-off when it comes to these  $n$ -steps methods. A lower  $n$  would correspond to a higher bias but low variance whereas a higher  $n$  would have the reversed effects. Why we see so little variation in terms of rate of convergence among the three different  $n$  could be that the trade-offs are pretty linear. Maybe it has something to do with the randomness of how the monster moves.

2)

The regular Q-learning method potentially suffers from the maximization bias (p. 134 in the course book), which means that we expect it to overestimate its action values early on in the training. The Double Q-learning method is the remedy to this bias and as such, it produces more reliable action values throughout the whole training. As we, in Figure 1, are looking at the actual rewards that result from the potentially biased policy and not the action-values themselves, the bias isn't apparent. We do, however, see a more rapid increase in reward early on from the Q-learning method, before the methods seem to converge at very similar values. This indicates that in this case, the maximization bias might actually be beneficial to early rewards.

3)

Convergence should be quicker as there are less paths to consider. Instead of  $\approx 4$  monster moves per agent action, with a deterministically moving monster the number of paths should at least be halved. Thus leading to a smaller state-action space, thus leading to faster exploring of the entire space, thus leading to faster convergence.

## Appendix 1: Python code

```
1 # %%
2 import numpy as np
3 import random
4 from matplotlib import pyplot as plt
5 import pandas as pd
6
7 # Gridworld parameters
8 N = 5 # grid size (N x N)
9 T = 30 # maximum time steps per episode
10 alpha = 0.1 # learning rate
11 gamma = 0.9 # discount factor
12 episodes = 10**5 # number of episodes per method
13
14 # Create grid coordinates
15 grid = [(r, c) for r in range(N) for c in range(N)]
16
17 # Define actions and their effects (using (row, col) convention)
18 actions = ["up", "right", "down", "left"]
19 actions_dict = {
20     "up": (-1, 0),
21     "right": (0, 1),
22     "down": (1, 0),
23     "left": (0, -1)
24 }
25
26 # Define rewards
27 rewards_dict = {
28     "collect_apple": 1,
29     "caught_by_monster": -1,
30     "empty": 0
31 }
32
33 # Parameters for epsilon decay
34 epsilon_0 = 1.0 # Initial exploration rate
35 epsilon_min = 0.01 # Minimum exploration
36 decay_rate = 0.0001 # Controls speed of decay
37
38 def get_epsilon(episode):
39     """ Exponentially decaying epsilon """
40     return epsilon_min + (epsilon_0 - epsilon_min) * np.exp(-decay_rate * episode)
41
42 def initial_positions():
43     """Return distinct starting positions for agent, monster, and apple."""
44     return tuple(random.sample(grid, 3))
45
46 def respawn_apple(agent_pos, monster_pos):
47     """Return a new apple position that is not occupied by agent or monster."""
48     available_positions = [pos for pos in grid if pos != agent_pos and pos != monster_pos]
49     return random.choice(available_positions)
50
51 def move(pos, action):
52     """
53     Given a position and an action, return the new position.
54     If the move is out-of-bounds, return the original position.
55     """
```

```

56     delta = actions_dict[action]
57     new_r = pos[0] + delta[0]
58     new_c = pos[1] + delta[1]
59     if 0 <= new_r < N and 0 <= new_c < N:
60         return (new_r, new_c)
61     else:
62         return pos
63
64 def step(state, agent_action, monster_action):
65     """
66     Execute one simultaneous step for agent and monster.
67     Returns next_state, reward, and a done flag.
68     """
69     agent_pos, monster_pos, apple_pos = state
70     new_agent_pos = move(agent_pos, agent_action)
71     new_monster_pos = move(monster_pos, monster_action)
72
73     # Check if agent and monster collide.
74     if new_agent_pos == new_monster_pos:
75         return (new_agent_pos, new_monster_pos, apple_pos), rewards_dict["caught_by_monster"], True
76
77     # Check for apple collection.
78     reward = rewards_dict["empty"]
79     if new_agent_pos == apple_pos:
80         reward = rewards_dict["collect_apple"]
81         new_apple_pos = respawn_apple(new_agent_pos, new_monster_pos)
82     else:
83         new_apple_pos = apple_pos
84
85     next_state = (new_agent_pos, new_monster_pos, new_apple_pos)
86     return next_state, reward, False
87
88 def epsilon_greedy(Q, state, epsilon):
89     """
90     Return an action chosen by the epsilon-greedy policy based on Q.
91     """
92     if state not in Q or random.random() < epsilon:
93         return random.choice(actions)
94     max_val = max(Q[state].values())
95     best_actions = [a for a, v in Q[state].items() if v == max_val]
96     return random.choice(best_actions)
97
98 def epsilon_greedy_double(Q1, Q2, state, epsilon):
99     """
100     For double Q-learning: choose an action using the sum of Q1 and Q2 values.
101     """
102     if state not in Q1 or state not in Q2 or random.random() < epsilon:
103         return random.choice(actions)
104     combined = {a: Q1[state][a] + Q2[state][a] for a in actions}
105     max_val = max(combined.values())
106     best_actions = [a for a, v in combined.items() if v == max_val]
107     return random.choice(best_actions)
108
109 # %%
110 # -----
111 # Off-Policy Q-Learning
112 # -----

```

```

113 Q_learning = {}
114 Q_learning_rewards = []
115
116 print("Starting Off-Policy Q-Learning...")
117 for episode in range(episodes):
118     state = initial_positions() # (agent_pos, monster_pos, apple_pos)
119     if state not in Q_learning:
120         Q_learning[state] = {a: 0.0 for a in actions}
121
122     epsilon = get_epsilon(episode)
123     total_reward = 0
124     t = 0
125     done = False
126     while t < T and not done:
127         # Select action using Q_learning's own epsilon-greedy.
128         agent_action = epsilon_greedy(Q_learning, state, epsilon)
129         monster_action = random.choice(actions)
130
131         next_state, reward, done = step(state, agent_action, monster_action)
132         total_reward += reward
133
134         if not done:
135             if next_state not in Q_learning:
136                 Q_learning[next_state] = {a: 0.0 for a in actions}
137                 best_next = max(Q_learning[next_state].values())
138             else:
139                 best_next = 0
140
141             Q_learning[state][agent_action] += alpha * (reward + gamma * best_next - Q_
learning[state][agent_action])
142             state = next_state
143             t += 1
144
145         # save accumulated episode reward
146         Q_learning_rewards.append(total_reward)
147
148         if (episode + 1) % 100 == 0:
149             print(f"Off-Policy Q-Learning Episode {episode+1}: Total Reward = {total_
reward}")
150
151 # %%
152 # -----
153 # 1-Step SARSA
154 # -----
155 Q_sarsa = {}
156 Q_sarsa_rewards = []
157
158 print("\nStarting 1-Step SARSA...")
159 for episode in range(episodes):
160     state = initial_positions()
161     if state not in Q_sarsa:
162         Q_sarsa[state] = {a: 0.0 for a in actions}
163     # Choose initial action using Q_sarsa.
164     agent_action = epsilon_greedy(Q_sarsa, state, epsilon)
165
166     epsilon = get_epsilon(episode)
167     total_reward = 0
168     t = 0

```

```

169 done = False
170 while t < T and not done:
171     if state not in Q_sarsa:
172         Q_sarsa[state] = {a: 0.0 for a in actions}
173     monster_action = random.choice(actions)
174     next_state, reward, done = step(state, agent_action, monster_action)
175     total_reward += reward
176
177     if not done:
178         if next_state not in Q_sarsa:
179             Q_sarsa[next_state] = {a: 0.0 for a in actions}
180             next_action = epsilon_greedy(Q_sarsa, next_state, epsilon)
181             Q_sarsa[state][agent_action] += alpha * (reward + gamma * Q_sarsa[next_
state][next_action] - Q_sarsa[state][agent_action])
182         else:
183             Q_sarsa[state][agent_action] += alpha * (reward - Q_sarsa[state][agent_
action])
184
185     state = next_state
186     if not done:
187         agent_action = next_action
188     t += 1
189
190     Q_sarsa_rewards.append(total_reward)
191     if (episode + 1) % 100 == 0:
192         print(f"SARSA Episode {episode+1}: Total Reward = {total_reward}")
193
194 # -----
195 # Double Q-Learning
196 # -----
197 # In double Q-learning, the behavior policy is derived from the sum of Q_double1 and Q
_double2.
198 Q_double1 = {}
199 Q_double2 = {}
200 Q_double_rewards = []
201
202 print("\nStarting Double Q-Learning...")
203 for episode in range(episodes):
204     state = initial_positions()
205     for Q in (Q_double1, Q_double2):
206         if state not in Q:
207             Q[state] = {a: 0.0 for a in actions}
208
209     epsilon = get_epsilon(episode)
210     total_reward = 0
211     t = 0
212     done = False
213     # Choose initial action using the double Q behavior policy.
214     agent_action = epsilon_greedy_double(Q_double1, Q_double2, state, epsilon)
215
216     while t < T and not done:
217         for Q in (Q_double1, Q_double2):
218             if state not in Q:
219                 Q[state] = {a: 0.0 for a in actions}
220             monster_action = random.choice(actions)
221             next_state, reward, done = step(state, agent_action, monster_action)
222             total_reward += reward
223

```



```

224         for Q in (Q_double1, Q_double2):
225             if next_state not in Q:
226                 Q[next_state] = {a: 0.0 for a in actions}
227
228         if not done:
229             next_action = epsilon_greedy_double(Q_double1, Q_double2, next_state,
epsilon)
230             # Randomly update one of the two Q-tables.
231             if random.random() < 0.5:
232                 best_action = max(Q_double1[next_state], key=Q_double1[next_state].get
)
233                 target = reward + gamma * Q_double2[next_state][best_action]
234                 Q_double1[state][agent_action] += alpha * (target - Q_double1[state][
agent_action])
235             else:
236                 best_action = max(Q_double2[next_state], key=Q_double2[next_state].get
)
237                 target = reward + gamma * Q_double1[next_state][best_action]
238                 Q_double2[state][agent_action] += alpha * (target - Q_double2[state][
agent_action])
239             else:
240                 # Terminal state update.
241                 if random.random() < 0.5:
242                     Q_double1[state][agent_action] += alpha * (reward - Q_double1[state][
agent_action])
243                 else:
244                     Q_double2[state][agent_action] += alpha * (reward - Q_double2[state][
agent_action])
245
246             state = next_state
247             if not done:
248                 agent_action = next_action
249             t += 1
250
251         Q_double_rewards.append(total_reward)
252         if (episode + 1) % 100 == 0:
253             print(f"Double Q-Learning Episode {episode+1}: Total Reward = {total_reward}")
254
255 # At the end, Q_learning, Q_sarsa, and (Q_double1, Q_double2) hold the learned action
values for each method.
256
257 # %%
258 # -----
259 # n-step SARSA Function
260 # -----
261 def n_step_sarsa(n, episodes):
262     """
263     Runs n-step SARSA on the gridworld for a given number of episodes.
264     No pre-initialization of all states is needed; states are added as encountered.
265
266     Parameters:
267         n : number of steps for bootstrapping
268         episodes : number of episodes to run
269
270     Returns:
271         Q_n : learned Q-value table (dictionary)
272         rewards_n: list of total rewards per episode
273     """

```

```

274 Q_n = {}
275 rewards_n = []
276 for episode in range(epochs):
277     epsilon_val = get_epsilon(episode)
278     # Initialize starting state and Q-values for that state if needed.
279     state = initial_positions()
280     if state not in Q_n:
281         Q_n[state] = {a: 0.0 for a in actions}
282     # Choose initial action using the epsilon-greedy policy.
283     action = epsilon_greedy(Q_n, state, epsilon_val)
284
285     # Lists to store the trajectory:
286     states = [state]          # states[0] is the initial state
287     actions_list = [action]   # actions_list[0] is the initial action
288     rewards_list = [0]       # rewards_list[0] is a dummy reward
289
290     T_episode = 30
291     t = 0
292     total_reward = 0
293
294     while True:
295         if t < T_episode:
296             monster_action = random.choice(actions)
297             next_state, reward, done = step(state, action, monster_action)
298             total_reward += reward
299             rewards_list.append(reward)
300             states.append(next_state)
301             if done:
302                 T_episode = t + 1
303             else:
304                 if next_state not in Q_n:
305                     Q_n[next_state] = {a: 0.0 for a in actions}
306                 next_action = epsilon_greedy(Q_n, next_state, epsilon_val)
307                 actions_list.append(next_action)
308             tau = t - n + 1
309             if tau >= 0:
310                 # Determine the upper index for the reward sum.
311                 limit = min(tau + n, T_episode)
312                 G = 0.0
313                 for i in range(tau + 1, limit + 1):
314                     G += (gamma ** (i - tau - 1)) * rewards_list[i]
315                 if tau + n < T_episode:
316                     G += (gamma ** n) * Q_n[states[tau + n]][actions_list[tau + n]]
317                 Q_n[states[tau]][actions_list[tau]] += alpha * (G - Q_n[states[tau]][
actions_list[tau]])
318                 t += 1
319                 if tau == T_episode - 1:
320                     break
321             # Update state and action only if available.
322             if t < len(states):
323                 state = states[t]
324             if t < len(actions_list):
325                 action = actions_list[t]
326             rewards_n.append(total_reward)
327             if (episode + 1) % 1000 == 0:
328                 print(f"n-step SARSA (n={n}) Episode {episode+1}: Total Reward = {total_
reward}")
329         return Q_n, rewards_n

```

```

330
331 # %% [code]
332 # -----
333 # Run n-step SARSA for different n values
334 # -----
335 # print("\nStarting n-step SARSA (n=2)...")
336 _, rewards_n2 = n_step_sarsa(2, episodes)
337
338 # print("\nStarting n-step SARSA (n=4)...")
339 _, rewards_n4 = n_step_sarsa(4, episodes)
340
341 # For comparison, n=1-step SARSA is equivalent to standard 1-step SARSA:
342 print("\nStarting n-step SARSA (n=1) [equivalent to 1-step SARSA]...")
343 _, rewards_n1 = n_step_sarsa(1, episodes)
344
345
346 # %%
347 # -----
348 # Plotting the Learning Curves
349 # -----
350 avg_len = 1000
351
352 # creating lists of average over 1000 last episodes
353 Q_sarsa_avg = pd.Series(Q_sarsa_rewards).rolling(window=avg_len, min_periods=1).mean().to_numpy()
354 Q_learning_avg = pd.Series(Q_learning_rewards).rolling(window=avg_len, min_periods=1).mean().to_numpy()
355 Q_dbl_avg = pd.Series(Q_double_rewards).rolling(window=avg_len, min_periods=1).mean().to_numpy()
356
357 # Compute rolling averages over the last 1000 episodes
358 n1_avg = pd.Series(rewards_n1).rolling(window=avg_len, min_periods=1).mean().to_numpy()
359 n2_avg = pd.Series(rewards_n2).rolling(window=avg_len, min_periods=1).mean().to_numpy()
360 n4_avg = pd.Series(rewards_n4).rolling(window=avg_len, min_periods=1).mean().to_numpy()
361
362 episodes_range = list(range(1, episodes + 1))
363 plt.figure(figsize=(10, 6))
364 plt.plot(episodes_range, Q_learning_avg, label="Off-Policy Q-Learning")
365 # plt.plot(episodes_range, Q_sarsa_avg, label="1-Step SARSA")
366 plt.plot(episodes_range, Q_dbl_avg, label="Double Q-Learning")
367 plt.plot(episodes_range, n1_avg, label="n-step SARSA (n=1)")
368 plt.plot(episodes_range, n2_avg, label="n-step SARSA (n=2)")
369 plt.plot(episodes_range, n4_avg, label="n-step SARSA (n=4)")
370 plt.xlabel("Episode")
371 plt.ylabel("Rolling average reward (1000 previous episodes)")
372 plt.legend()
373 plt.grid(True)
374 plt.show()
375 # %%

```

## References

Sutton, Richard S (2018). “Reinforcement learning: An introduction”. In: *A Bradford Book*.