

Reinforcement learning: Project 4

Florence Hugh, August Jonasson, Amanda Tisell

March 21, 2025

Introduction

In this project our goal is to train an agent to play the game Othello, which is a modern version of Reversi¹. Othello is a turn-based strategy board game for two players played on an 8×8 grid. Players take turns placing one piece on an empty spot on the board with their assigned color. After a play is made, the pieces of the opponent's color that lie bounded by the current player's color are turned over. The game ends when the board is filled or when none of the players can make a move. The player who controls the most pieces at the end wins the game. Figure 1 shows an example of the first few moves in a game.




Figure 1: Starting state and first two moves of a game.

Some challenges and how we address them

Since Othello is a two-player game, we need to consider the agent/opponent setting. A natural idea would be to have one set of values for black and a completely different set of values for white, thus yielding two separate policies. When dealing with a two-player setting such as this, however, it turns out that we can avoid having to train two agents simultaneously by using self-play in combination with a minimax algorithm, thus allowing for both the black and white pieces to play against the same set of values.

Othello, similarly to Chess and AlphaGo, is unsolved and has a huge state space (around 10^{28} possible board positions on an 8×8 grid). Handling this size would require a large amount of compute which we do not have. Therefore, we decide to cut down the size of the board to a 6×6 grid instead. This results in a state space of around 10^{12} , which is still large enough to prevent accidentally brute-forcing the problem, but much more manageable within the scope of this project. A possibly even more important

¹<https://en.wikipedia.org/wiki/Reversi>

advantage of this simplification is that 6×6 Othello is solved, with the white pieces having a clear advantage². This means that it will now be much easier to assess the results of the training, since we now know what behavior to look for. Ideally, we would like to see the agent converge towards white winning every game.

As in most reinforcement learning systems, we have a setting of sparse rewards (only at the end of games) and the issue of balancing exploration/exploitation properly. We want to be able to learn meaningful patterns in a reasonable amount of time without having to go through the whole decision tree of possible games. Our attempt at addressing both of these problems (which is heavily inspired by the AlphaGo article Silver et al. 2016) is to use a Monte Carlo tree search algorithm in combination with a neural network-approximated policy that, simply put, favors deep over wide exploration such that rewards become more frequent, thus speeding up the learning.

MDP framing

The environment is a 6×6 grid where the dynamics are completely deterministic. The number of unique board positions in 6×6 Othello has been approximated to 10^{12} ($3^{32} \cdot 2^4 \approx 10^{16}$ and then divided by some number that accounts for what proportion of these actually represent legal board positions). A state will be represented as a 6×6 matrix with entries being

$$s_{ij} = \begin{cases} -1 & \text{if square } (i, j) \text{ is occupied by a white piece,} \\ 0 & \text{if square } (i, j) \text{ is not occupied by any piece,} \\ 1 & \text{if square } (i, j) \text{ is occupied by a black piece.} \end{cases} \quad (1)$$

The action space of size 32 includes all squares on the board, excluding the four starting pieces. However, most of these actions will be illegal. According to a paper that claims to have weakly solved Othello, there are on average 10 legal actions per move on a 8×8 grid (Takizawa 2023). This leads us to believe that a 6×6 grid may have on average 5-8 legal actions from a given state. An action a is represented as a tuple (i, j) , which is the position where the piece is placed on the board.

A reward is only given at the end of each game and will be defined as

$$R(s, a) = \begin{cases} 1 & \text{if black wins the game} \\ -1 & \text{if white wins the game} \\ 0 & \text{if the game is drawn} \end{cases} \quad (2)$$

²https://en.wikipedia.org/wiki/Computer_Othello

Monte Carlo tree search (MCTS)

Monte Carlo tree search (Sutton, Barto, et al. 2018, ch. 8.11) is a type of decision-time planning algorithm based on Monte Carlo control applied to simulations starting from the root state, that is, a type of rollout algorithm.

Decision-time planning (Sutton, Barto, et al. 2018, ch. 8.8) is a type of planning that focuses on a particular state. That is, instead of simulating experience for the whole state space, it only plans ahead from the current state s_t . Using simulated experience to gradually improve a policy or value function or to select an action for the current state. The values and policy are specified to the current state and the action choices available there. In decision-time planning there is a tradeoff between better policies and the time that is needed to simulate enough trajectories to obtain good value estimates.

A rollout algorithm (Sutton, Barto, et al. 2018, ch. 8.10) is a decision-time planning algorithm based on Monte Carlo control applied to simulated trajectories that all begin in the current state. Each rollout produce Monte Carlo estimates of action-values for each current state and for a given policy, usually called the rollout policy. The aim of a rollout algorithm is to improve the rollout policy and not find the optimal policy like in regular Monte Carlo methods.

In MCTS, as in a rollout algorithm, each execution is an iterative process that simulates many trajectories, each trajectory starting from the current state s_t , running to a terminal state s_T . This iterative process consists of four steps; selection, expansion, rollout (simulation) and backpropagation.

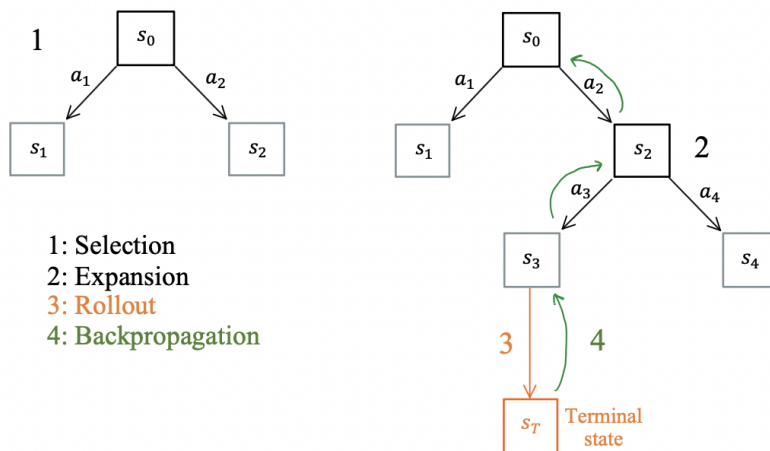


Figure 2: Monte Carlo tree search

These four steps continues to execute, starting each time at the tree's root node (initial state), until some criterion is fulfilled. This criterion can for instance be iterating for a fixed number of iterations (fixed number of terminal states reached with the tree policy).

1. Selection

In the selection part, a node (state) is chosen according to a tree policy. The tree policy balances exploration and exploitation, where it's action selection can for instance be based on ϵ -greedy or an

upper confidence bound (UCB) selection formula

$$a_t = \arg \max_a \left(Q(s, a) + c \sqrt{\frac{\log N(s)}{N(s, a)}} \right). \quad (3)$$

The $Q(s, a)$ is an average action-value (updated by the backpropagation, explained in part 4), c is a constant that controls the greediness of the action selection, $N(s)$ is the number of visits in current state s and $N(s, a)$ is the number of visits in the next state given action a .

The MCTS traverses down the decision tree and whenever a state is visited, the next action is selected among the possible actions from this current state according to equation 3. This process continues until a leaf node (unexplored and not terminal state) is reached in which an expansion is applied.

2. Expansion

The decision tree is expanded from the current state by connecting all possible actions to their respectively next state. From these new states one is selected based on the tree policy used in the selection step. However, all these newly added states will have a $N(s, a)$ count equal to zero. Since the MCTS prioritizes untried actions, one of these states will simply be picked randomly as a zero count will be seen as a a_t value of $\pm\infty$. Whenever a zero counted state-action is chosen we proceed to the next step, the rollout.

3. Rollout

During the rollout, a simulation is applied from the current state. This simulates a whole episode starting from the current state until a terminal state is reached. The actions selected during the simulation are generated using a rollout policy and does not necessarily follow the chosen tree policy. The rollout policy is usually a simpler policy allowing for faster computation. It could for example follow a policy where actions are chosen uniformly. When the terminal state is reached, a rollout reward R_r is received. The reward will be backpropagated and finally, all values generated by the rollout will be discarded.

4. Backpropagation

In the final phase, called the backpropagation, the reward observed from doing the rollout is backpropagated through the tree, updating all state-action values from where we started the rollout (the leaf) up to the root. The update rule is the in-place update of the average, given by

$$Q_{n+1}(s, a) \leftarrow Q_n(s, a) + \frac{1}{N(s, a)} [R_r - Q_n(s, a)], \quad (4)$$

(Sutton, Barto, et al. 2018, eq. (2.5)) where $Q_n(s, a)$ is the old state-action-value, $N(s, a)$ is the number of visits to state-action pair s, a , R_r is the reward yielded from the rollout, and $Q_{n+1}(s, a)$ is the updated value.

Modifying MCTS

In order to speed up the tree traversal and gain better control of the exploration (such that rewards become more frequent) we now make some modifications to regular MCTS.

First, to address the agent/opponent setting of the system, we solve this problem by using a minimax³ algorithm in combination with self-play. As such, we reframe the system as an agent learning how to play Othello against itself, rather than an agent playing Othello against an opponent. The minimax algorithm is a decision rule strategy used in game theory in order to find optimal moves in zero-sum games (when a player's gain is balanced out by the other player's loss). The algorithm assumes that one player tries to maximize its score, while the opponent tries to minimize it. This fits well into our model as our reward structure results in a zero-sum game.

We have incorporated this algorithm into the Monte Carlo tree structure by adding additional information about which color's move it is at each state. By doing this the MCTS has an indicator of when to minimize/maximize the next state-action value. The action selection part of the tree will handle the minimax situation by using formula

$$a_t = \underset{a}{\overset{(\arg \min)}{\arg \max}} \left(Q(s, a) \pm c \frac{P(s, a) \sqrt{N(s)}}{N(s, a)} \right), \quad (5)$$

where the black (white) player will try to maximize (minimize) the return.

In addition, a policy $P(s, a)$ is added in order to make the agent more selective toward states that have already seen a lot of visits, thus pushing the actual exploration down along a few branches instead of higher up among many branches. This will lead to rewards being less sparse in the training, and as such, will speed up learning. The policy itself is retrieved through a neural network, which we will present in its own section.

Lastly, the rollout policy is replaced by function approximation of state-action values, as constantly letting the games play out to their end is too slow. Now, instead of playing them out, we simply input the leaf state into a function and get an approximation of the value back immediately. This function approximation is also handled by a neural network which we will present in its entirety in the next section. The backpropagation still remains as before. The only difference now is that instead of the actual reward, it is the approximated value that gets backpropagated, unless we have actually reached a terminal state - then the real reward is backpropagated.




Figure 3: Modified Monte Carlo tree search

³<https://en.wikipedia.org/wiki/Minimax#Minimax>

Value approximation network

The function approximation of the action-values is handled by a deep convolutional neural network. In this section we will briefly present the architecture of the network and also explain how, and on what data, the initial training was done. For specific hyperparameter settings (such as batch-sizes, learning rates, etc.) we refer to the code appendix (line 749 in the code).

The architecture of the value neural network is heavily inspired by resnet-18 (He et al. 2016). Figure 4 depicts the network in its entirety.

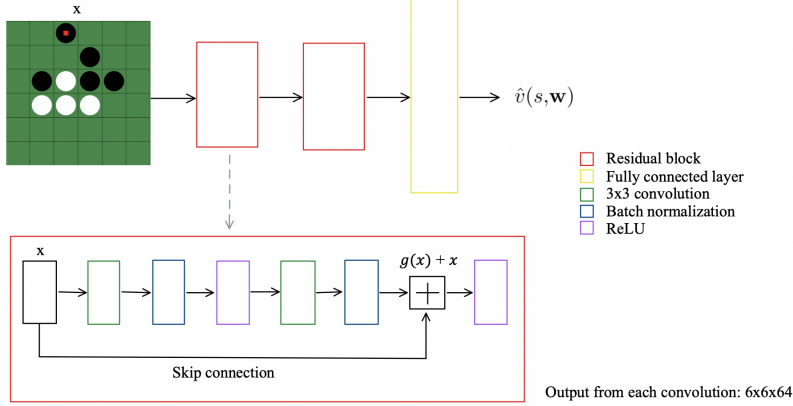



Figure 4: Value network architecture.

The input is a 6×6 matrix with elements according to (1) in the MDP framing, the targets are the Monte Carlo returns associated with the game from which that particular board position was retrieved, i.e. $-1, 0$ or 1 , the output is the full-gradient action-value approximation, and finally, the loss function is the Mean Squared Value Error with $\mu(s)$ set to constant $1/|\text{data}|$ (Sutton, Barto, et al. 2018, p. 199).

The data used for the initial training of the value network consists of unique board positions retrieved from one hundred thousand randomly simulated games, with each one labeled according to the reward that was given at the end of its respective game. In total, this yields a dataset of size $\approx 2 \cdot 10^6$, which is then divided into a training dataset and a validation dataset according to a 80-20 split. Figure 5 depicts the training curve for 150 epochs. From the validation curve, we see no indication of overfitting, and the training curve seems to converge. As such, we deduce that the training is healthy. Looking at the actual loss, however, which is squared, we notice that for returns in the range $[-1, 1]$ it seems relatively large. This might be an indication that our network is not deep enough to successfully capture the complex relationships between board positions and returns. Given more time and resources, adding more residual blocks is a potential improvement for the project. See code appendix for the full `tensorflow` implementation.


Figure 5: Initial training of the value network.



Policy approximation network

The policy approximation is handled by a plain, fully connected, feedforward neural network with linear inputs up until and including the output layer, which uses a softmax function in order to model the probability distribution over the legal actions in the given state. The architecture is summarized in Figure 6. Again, refer to the code appendix for specific hyperparameter settings (line 802 in the code).


Since this network only requires the states as a sort of identifier to know which actions are legal, the board position matrices can be flattened to a 36-dimensional vector before being input into the network, thus avoiding the need for convolutions to preserve more complex relations. Most likely, the states can be mapped to an even lower dimension and still act as unique identifiers. This is a potential improvement for the network, in terms of lowering the necessary complexity. The targets here are the normalized number of times each action had been taken from the given state in that particular tree, thus forming an approximate probability distribution over the actions. Since the number of actions can vary depending on the state, the output dimension is kept at a constant 36. This means that the target will be of this dimension as well, but with zeros on all positions that are not legal actions from the given state.



This means that the predicted probabilities will not form a perfect distribution over the legal actions, but the hope is that the bulk of the probability mass will still fall here, as the remaining actions are fitted to zeros. Dealing with this is another potential improvement. Finally, the Kullback-Leibler divergence⁴ was used as loss function (minimization of distance between probability distributions).

The dataset used for the initial training of the policy network is a set of one thousand games simulated using the regular *upper-confidence bound* formula (eq. (3)) but with the value network instead of the rollout policy, thus forming a sort of intermediary step before transitioning completely to the model in Figure 3. Ideally, this is where we would want to use prerecorded “expert” experience, but due to not being able find any such databases for 6×6 Othello, we settle for this solution, and assume that the MCTS with the value network alone is sufficient. The training can be seen in Figure 7 and looks healthy. Since the targets represent unique state-action visits, we are ensured that there are no identical matches between the training and validation set. This indicates that the model generalizes well to unseen data.

Figure 7: Initial training of the value network.



⁴https://en.wikipedia.org/wiki/Kullback%E2%80%93Leibler_divergence

Training cycle

With the modified MCTS, the value network and the policy network, we are now ready to define a *training cycle*, which is shown in Figure 8. Using the initially trained neural networks, we begin with simulating games from the Monte Carlo tree with a predefined number of episodes before termination. These games are then stored in a replay buffer which upon termination of the tree is sent into training of the neural networks for a number of epochs. Finally the updated networks are sent back to the tree, where a new step of the cycle can begin.




Figure 8: The agent training cycle.

In the actual implementation of the model, we let one tree iterate until it had reached terminal states 25 times, for 30 cycles. Even a single cycle often took as long as 30 minutes to complete due to the tree exploring along several different branches before reaching terminal states. In total, with the initial training of the value and policy networks included, approximately 15 hours was spent on training the model (on a CPU).

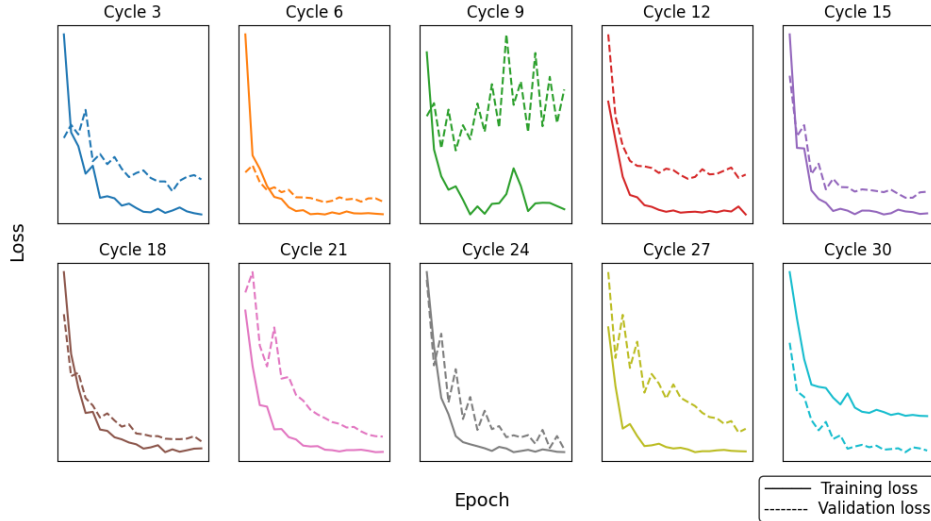


Figure 9: Value network training from ten selected cycles.

Figures 9 and 10 depict the training curves of the neural networks within the 30 training cycles. Overall the training looks very healthy in both networks. Cycle 9 within the value network has a validation




Figure 10: Policy network training from ten selected cycles.

curve that indicates a poor fit, which would be interesting to examine further. Cycle 30 from the value network has a validation curve that lies entirely below the training curve. This is actually because the training has converged such that the training and validation datasets are perfectly correlated. The regularization in the training is what causes the validation curve to fall below. In the policy network training we don't see any anomalies, and we can now move on to the results of the agent.

Results

Figure 11 shows the result after training the model on 30 training cycles with 25 episodes each. As seen in Figure 11, the white piece player becomes increasingly more dominant and, at the end, completely dominating. This is a promising result as 6×6 Othello is solved with the result that white has an advantage.




Figure 11: Reward distribution of each training cycle.




Figure 12: Terminal state distribution of cycle 5 (left) and cycle 30 (right).

Delving into some of the specific cycles now, we can see from Figure 12, for example, that even though white wins the first 6 episodes in the cycle, the fact that it wins in 6 different ways indicates that black is fighting back. Eventually, black even manages to get back with 6 wins in the same terminal state. Why white allows for this same terminal state to be reached so many times in a row could have something to do with the fact that white previously has gotten high rewards along this branch and therefore requires a couple of iterations before it realizes that better options exist.

Finally, though, which can be seen from Figure 12, white takes over completely and manages to force the game towards a single terminal state each game in the whole cycle. This happened in four other cycles towards the end as well, further strengthening the hypothesis that our training converged.

Discussion

The final result indicates that our agent has trained well. However, to gain more insight into the agents behavior, a deeper analysis of the trained decision tree is required. This involves examining the structure of the tree to find meaningful patterns, such as frequently visited paths and areas where the agent shows uncertainty or bias, and map out specific trajectories and compare them to solutions of 6×6 Othello. Since the game is solved, it would actually be possible to see if an optimal policy has been reached.

We mention that initially we worked on a 4×4 grid where a large amount of time was spent on implementing the methods. This turned out to be too simple as we ended up brute-forcing the whole tree, resulting in an optimal solution during the first cycle. Then we jumped straight to the 8×8 grid in which we realized we did not have enough computational power. Since our code was written for the general case, we seamlessly ended up with the 6×6 grid eventually.

Possible Improvements

Neural Networks

The policy network needs further improvements seeing how the probability mass of the output did not match well with the legal actions deeper down the tree. Meaning that in some cases, as little as 50 % of the probability mass was distributed to the legal actions. One way to solve it could be by adding a masking part which zeros out the illegal actions and weighted up the legal actions so they sum to one. There is probably a better way to improve this but we would need more time to figure it out.

Another improvement could be to simply train one neural network instead of having two separate ones. This network would then have as output both the actions and the values. By only training one network, we may speed up the training cycle slightly. It was, however, the tree traversal that ate up most of the training time, due to the size of the state space. So, any gains in the traversal would be even more valuable.

Parallel Processing

To train the tree, it needed to be traversed many times. Instead of just using one agent to traverse down the tree, we could have trained several agents simultaneously to enhance the process. There should also be a way for parallelizing the networks and the tree.

Handle Symmetry

We know for sure that the first four positions in which black can put its pieces are completely symmetric. With this knowledge, we could have fixed the first black piece to one position and then only trained that part of the tree to optimize its performance on one branch. Then the result in this branch could somehow map the other starting positions to equivalent actions.

Play against the agent

Ideally, it would be nice to be able to play with the agent. This way we could see if it actually wins all the games (given that it plays the white pieces) when playing against a human. We started building a graphical user interface (GUI) to play against the agent. However, because our agent had not explored enough state-action pairs, it could not make an action fast enough as it had to explore the tree from this unseen state. A way to solve this could be to either build a larger tree with the already trained weights from the networks (but this takes quite a time) or we could have added a time

horizon threshold such that the agent only had a set amount of time to think. As it was implemented now, it spent a long time thinking (planning ahead) on each move which made the games very long. We are unfortunately not entirely sure how well it performs against a human, as of now.

References

- He, Kaiming et al. (2016). “Deep residual learning for image recognition”. In: *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778.
- Silver, David et al. (2016). “Mastering the game of Go with deep neural networks and tree search”. In: *nature* 529.7587, pp. 484–489.
- Sutton, Richard S, Andrew G Barto, et al. (2018). *Reinforcement learning: An introduction*. Vol. 1. 1. MIT press Cambridge.
- Takizawa, Hiroki (2023). “Othello is solved”. In: *arXiv preprint arXiv:2310.19387*.

Appendix A: Python code

```
1 import copy
2 import numpy as np
3 import pickle
4 import random
5 import tensorflow as tf
6 from MCT_Othello_classes import *
7 from sklearn.model_selection import train_test_split
8 from matplotlib import pyplot as plt
9 import matplotlib.ticker as mticker
10 import keras
11 from keras import layers, regularizers, Model
12 from othello_rl_helper_fcts import *
13
14 # 0=blank, 1=tot_black, -1=white
15 class OthelloBoard():
16     dirx = [-1, 0, 1, -1, 1, -1, 0, 1]
17     diry = [-1, -1, -1, 0, 0, 1, 1, 1]
18
19     def __init__(self, n):
20         self.n = n
21         self.board = [[0 for _ in range(n)] for _ in range(n)]
22         self.to_play = 1 #keep track of players turn 1=tot_black -1=white
23         # self.pass_counter = 0
24         self.reset_game()
25
26     def reset_game(self):
27         n = self.n
28         # self.pass_counter = 0
29         self.to_play = 1
30         self.board = [[0 for _ in range(n)] for _ in range(n)]
31         board = self.board
32
33         # set up initial bricks
34         z = (n - 2) // 2
35         board[z][z] = -1
36         board[n - 1 - z][z] = 1
37         board[z][n - 1 - z] = 1
38         board[n - 1 - z][n - 1 - z] = -1
39
40         return board
41
42     def print_board(self):
43         n = self.n
44         board = self.board
45         m = len(str(n - 1))
46         for y in range(n):
47             row = ''
48             for x in range(n):
49                 row += str(board[y][x])
50                 row += ' ' * m
51             print(row + ' ' + str(y))
52         print("")
53         row = ''
54         for x in range(n):
55             row += str(x).zfill(m) + ' '
56         print(row + '\n')
57
58     def make_move(self, curr_state, action, to_play):
59         x = action[0]
60         y = action[1]
61         if self.check_valid_move(curr_state, x, y, to_play):
62             n = self.n
63             bricks_taken = 0 # total number of opponent pieces taken
```

```

64
65     curr_state[y][x] = to_play
66     for d in range(len(self.dirx)): # 8 directions
67         bricks = 0
68         for i in range(n):
69             dx = x + self.dirx[d] * (i + 1)
70             dy = y + self.diry[d] * (i + 1)
71             if dx < 0 or dx > n - 1 or dy < 0 or dy > n - 1:
72                 bricks = 0; break
73             elif curr_state[dy][dx] == to_play:
74                 break
75             elif curr_state[dy][dx] == 0:
76                 bricks = 0; break
77             else:
78                 bricks += 1
79         for i in range(bricks):
80             dx = x + self.dirx[d] * (i + 1)
81             dy = y + self.diry[d] * (i + 1)
82             curr_state[dy][dx] = to_play
83         bricks_taken += bricks
84     return (curr_state, bricks_taken)
85 else:
86     return print("Not valid move, retry")
87
88 def check_valid_move(self, curr_state, x, y, to_play):
89     """
90     Function checks playable moves. First if the agent is within the board, then
91     checks
92     if the spot is occupied by a tot_black or white brick and finally, if the
93     player do not
94     take any of the opponents bricks, then it is not a legal move.
95     """
96     if x < 0 or x > self.n - 1 or y < 0 or y > self.n - 1:
97         return False
98     if curr_state[y][x] != 0:
99         return False
100     (_, totctr) = self._check_valid_move(copy.deepcopy(curr_state), x, y, to_play)
101     if totctr == 0:
102         return False
103     return True
104
105 def _check_valid_move(self, board, x, y, to_play):
106     """
107     Helper function to check_valid_move function to not overwrite the playing
108     board if move is illegal
109     """
110     n = self.n
111     bricks_taken = 0 # total number of opponent pieces taken
112
113     board[y][x] = to_play
114     for d in range(len(self.dirx)): # 8 directions
115         bricks = 0
116         for i in range(n):
117             dx = x + self.dirx[d] * (i + 1)
118             dy = y + self.diry[d] * (i + 1)
119             if dx < 0 or dx > n - 1 or dy < 0 or dy > n - 1:
120                 bricks = 0; break
121             elif board[dy][dx] == to_play:
122                 break
123             elif board[dy][dx] == 0:
124                 bricks = 0; break
125             else:
126                 bricks += 1
127         for i in range(bricks):
128             dx = x + self.dirx[d] * (i + 1)

```

```

126         dy = y + self.diry[d] * (i + 1)
127         board[dy][dx] = to_play
128         bricks_taken += bricks
129         return (board, bricks_taken)
130
131     def move_generator(self, curr_state, to_play):
132         possibleMoves = []
133         for i in range(self.n):
134             for j in range(self.n):
135                 if (self.check_valid_move(curr_state, i, j, to_play)):
136                     possibleMoves.append((i, j))
137         return possibleMoves
138
139     def find_winner(self, curr_state):
140         tot_black = 0
141         tot_whites = 0
142
143         for i in range(self.n):
144             for j in range(self.n):
145                 if (curr_state[i][j] == -1):
146                     tot_whites += 1
147                 elif (curr_state[i][j] == 1):
148                     tot_black += 1
149
150         if (tot_black == tot_whites):
151             return 0
152         elif (tot_black > tot_whites):
153             return 1
154         else:
155             return -1
156
157
158
159     class MCTSNode(OthelloBoard):
160         def __init__(self, n, state, to_play, parent=None, parent_action=None):
161             super().__init__(n)
162             self.terminal_visits = 0
163             self.state = state
164             self.parent = parent
165             self.parent_action = parent_action
166             self.child_nodes = []
167             self._nof_visits = 0
168             self.player_turn = to_play
169             self.q_value = 0
170             self.p_action = 0
171             self.avg_q_value = 0
172             self.pass_counter = 0
173             self._untried_actions = self.untried_actions()
174             self.all_visited = False
175
176
177         def untried_actions(self):
178             self._untried_actions = self.move_generator(self.state, self.player_turn)
179
180             if len(self._untried_actions) == 0 and self.pass_counter != 2:
181                 self.pass_counter += 1
182                 self.player_turn *= -1
183                 self.untried_actions()
184
185             return self._untried_actions
186
187         def expand(self):
188             action = self._untried_actions.pop()
189             next_state, _ = self.make_move(copy.deepcopy(self.state), action, self.player_
turn)

```



```

190     next_player = self.player_turn*-1
191     child_node = MCTSNode(self.n, next_state, next_player, parent=self, parent_
action=action)
192     self.child_nodes.append(child_node)
193     if len(self._untried_actions) == 0:
194         self.all_visited = True
195     return child_node
196
197 def update_q(self, val):
198     self.q_value = val
199     # return self.q_value
200
201 # generalize this function such that it works for something
202 def uniform_policy(self):
203     """
204     Initializes a policy uniformly over all legal actions.
205     """
206     nof_actions = len(self.child_nodes)
207     return 1 / nof_actions
208
209 def backpropagate(self, q_NN):
210     # self.acum_q_value += q_NN
211     self._nof_visits += 1
212     self.avg_q_value += (q_NN - self.avg_q_value)/self._nof_visits #  $Q_{n+1}$ 
213     if self.parent: # Check if list is empty
214         self.parent.backpropagate(q_NN)
215
216 def best_child(self, c):
217     """
218     Minimax for training two agents
219     """
220     if self.player_turn == 1: # max
221         UCB_values = [child.avg_q_value + c * child.p_action * np.sqrt(self._nof_
visits) / child._nof_visits
222                        for child in self.child_nodes]
223         # + c_param*np.sqrt(np.log(self._nof_visits)/child._nof_visits)
224         return self.child_nodes[np.argmax(UCB_values)]
225     elif self.player_turn == -1: # min
226         UCB_values = [child.avg_q_value - c * child.p_action * np.sqrt(self._nof_
visits) / child._nof_visits
227                        for child in self.child_nodes]
228         # - c_param*np.sqrt(np.log(self._nof_visits)/child._nof_visits)
229         return self.child_nodes[np.argmin(UCB_values)]
230
231 def unpack_positions_returns(games):
232     """
233     Takes in games in the form of a list of tuples where the first element of the
234     tuple is a list of board positions representing one full game, and the second
235     element in the tuple is the reward (-1, 0, or 1) from that game.
236
237     Returns a list of tuples where each tuple consists of a single board position
238     and the reward associated with the game where that board position came from.
239     This is a (state s_t, return G_t) tuple in RL terms.
240     """
241     data = []
242     for game in games:
243         for position in game[0]:
244             data.append((position, game[1]))
245     return data
246
247 def unique_board_positions(state_return_tuples):
248     """
249     Takes in list of tuples (state, return) where state is a board position and
250     returns
251     a list of only the unique tuples.

```

```

251
252     This function will most likely be used only once, on the purely random simulated
253     data.
254     """
255     unique_dict = {}
256     for state, target in state_return_tuples:
257         # Convert the matrix to a hashable representation using .tobytes()
258         key = (state.tobytes(), target)
259
260         # Only add unique tuples
261         if key not in unique_dict:
262             unique_dict[key] = (state, target)
263
264     # Saving the unique (board, reward) tuples as data list
265     return list(unique_dict.values())
266
267 def value_predictors_targets(state_return_tuples):
268     """
269     Takes in list of tuples (state, return) where state is board position and return
270     is
271     observed reward from game.
272
273     And returns them in the form of X, y, ready to be used in a model.
274     """
275     # orders data into predictors and targets
276     X = np.array([x for x, _ in state_return_tuples])
277     y = np.array([y for _, y in state_return_tuples])
278
279     # expanding X to include #channels=1
280     X = np.expand_dims(X, axis=-1)
281
282     return X, y
283
284 # Unpickling and saving the games to a list
285 def read_files(path):
286     with open(path, "rb") as fp:
287         boards = pickle.load(fp)
288     return boards
289
290 def create_dataset(X, y, batch_size=128, shuffle_buffer_size=10000):
291     """
292     X and y should already be processed according to above preprocessing
293     functions before being passed to this function.
294
295     This function is simply to make training more efficient from memory.
296     """
297     # Create dataset from tensor slices
298     dataset = tf.data.Dataset.from_tensor_slices((X, y))
299
300     # Shuffle the dataset (reshuffles each epoch)
301     dataset = dataset.shuffle(buffer_size=shuffle_buffer_size, reshuffle_each_
302                             iteration=True)
303
304     # Batch the dataset
305     dataset = dataset.batch(batch_size)
306
307     # Prefetch for performance optimization
308     dataset = dataset.prefetch(tf.data.AUTOTUNE)
309
310     return dataset
311
312 def map_states_to_action_visits(sav):
313     state_to_action_visits = {}
314     for state, action, count in sav:
315         key = (state.tobytes())

```

```

313         if key not in state_to_action_visits:
314             state_to_action_visits[key] = [(action, count)]
315         else:
316             state_to_action_visits[key].append((action, count))
317
318     return state_to_action_visits
319
320 def map_actions_to_integers(n):
321     actions = [(i,j) for i in range(n) for j in range(n)]
322     map_actions = dict(zip(actions, [i for i in range(n*n)]))
323     return map_actions
324
325
326 def policy_predictors_targets(sav, n):
327     """
328     Takes in a list of tuples (state, action, count) where state is board position,
329     action is the action taken in this state and the count is the number of visits
330     to the next state when taking the action.
331
332     It should be the  $N(s,a)$  count from the UCB algorithm.
333
334     The predictors are the states and the target will be the  $N(s,a)$  count
335     """
336     map_sav = map_states_to_action_visits(sav)
337     mapped_actions = map_actions_to_integers(n)
338     X = []
339     y = []
340
341     for state, value in map_sav.items():
342         x = np.frombuffer(state, dtype=int)
343         X.append(x)
344         y_i = np.zeros(n*n)
345         total_visits = 0
346         for action, visits in value:
347             total_visits += visits
348             y_i[mapped_actions[action]] = visits
349
350         y_i = y_i/total_visits
351         y.append(y_i)
352
353     X = np.array(X)
354     y = np.array(y)
355
356     return X, y
357
358 def episode(node):
359     episode_list = [np.array(node.state)]
360     if node.parent:
361         episode_list.extend(episode(node.parent))
362     return episode_list
363
364 def treetraversal(node, res):
365     """
366     Recursive helper function to state_action_visits.
367     """
368     if not node:
369         return
370
371     if node.parent != None:
372         tup = (np.array(node.parent.state), node.parent_action, node._nof_visits)
373         res.append(tup)
374
375     for child in node.child_nodes:
376         treetraversal(child, res)
377

```

```

378
379 def state_action_visits(root):
380     """
381     Given MCTS, recursively goes through the tree and returns a list with tuples
382     containing state s, action a and N(s,a).
383     """
384     res = []
385     treetraversal(root, res)
386     return res
387
388
389 def results_distribution(episodes):
390     """
391     Given episodes, return a tuple with win, draw and loss.
392     """
393     win = 0
394     loss = 0
395     draw = 0
396
397     for i in range(len(episodes)):
398         if episodes[i][1] == -1:
399             loss += 1
400
401         elif episodes[i][1] == 0:
402             draw += 1
403
404         elif episodes[i][1] == 1:
405             win += 1
406
407     return (win, draw, loss)
408
409
410 def terminal_state_visits(episodes):
411     """
412     Given episodes, returns a dictionary with terminal state as key
413     and number of visits in that terminal state as value.
414     """
415     term_state_visits = {}
416
417     for i in range(len(episodes)):
418         term_state = (episodes[i][0][0].tobytes(),) # Make terminal state hashable
419         term_state_visits[term_state] = term_state_visits.get(term_state, 0) + 1
420
421     return term_state_visits
422
423
424 def bar_plot_term_states(episodes):
425     """
426     Given episodes, plot the distribution over how often a terminal state is visited.
427     """
428     win, draw, loss = results_distribution(episodes)
429     tot_games = len(episodes)
430     win_proc = np.round(win/tot_games*100, decimals=4)
431     draw_proc = np.round(draw/tot_games*100, decimals=4)
432     loss_proc = np.round(loss/tot_games*100, decimals=4)
433
434     term_state_visits = terminal_state_visits(episodes)
435     x_vals = [i for i in range(len(term_state_visits))]
436     plt.bar(x_vals, term_state_visits.values())
437
438     # Add a text box with game results
439     textstr = f"Total games = {tot_games}\nWin: {win_proc}%\nDraw: {draw_proc}%\nLoss: {loss_proc}%"
440
441     # Position the text box

```

```

442 props = dict(boxstyle='round', facecolor='white', alpha=0.8)
443 plt.text(0.68, 0.9, textstr, transform=plt.gca().transAxes, fontsize=9,
444         verticalalignment='top', bbox=props)
445
446 plt.xlabel("Terminal states")
447 plt.ylabel("Number of visits")
448 plt.grid()
449
450
451 def online_value_training(value_model, replay_buffer, epochs=20, batch_size=64):
452     """
453     Given a replay buffer in the form of list of tuples, (where each tuple consists of
454     a list of board positions from one game, and the reward from said game) and a
455     value model,
456     trains this model using the replay buffer through given number of epochs.
457
458     Returns the trained value_model
459     """
460     state_return_tuples = unpack_positions_returns(replay_buffer)
461     X_buffer, y_buffer = value_predictors_targets(state_return_tuples)
462     X_train_onl, X_test_onl, y_train_onl, y_test_onl = train_test_split(X_buffer, y_
463                                     buffer,
464                                     test_size=0.2,
465                                     random_state
466                                     =80085)
467
468     train_dataset = create_dataset(X_train_onl, y_train_onl)
469     val_dataset = create_dataset(X_test_onl, y_test_onl)
470
471     history_onl = value_model.fit(
472         train_dataset,
473         batch_size=batch_size,
474         epochs=epochs,
475         validation_data=val_dataset
476     )
477
478     return value_model, history_onl
479
480
481 def online_policy_training(policy_model, tree, epochs=20, n=6, batch_size=64):
482     """
483     Given a replay buffer in the form of list of state-actions and their visits and a
484     policy model,
485     trains this model using the replay buffer through given number of epochs.
486
487     Returns the trained policy_model
488     """
489     savs = state_action_visits(tree)
490     X_buffer, y_buffer = policy_predictors_targets(savs, n)
491     X_train_onl, X_test_onl, y_train_onl, y_test_onl = train_test_split(X_buffer, y_
492                                     buffer,
493                                     test_size=0.2,
494                                     random_state
495                                     =80085)
496
497     train_dataset = create_dataset(X_train_onl, y_train_onl)
498     val_dataset = create_dataset(X_test_onl, y_test_onl)
499
500     history_onl = policy_model.fit(
501         train_dataset,
502         batch_size=batch_size,
503         epochs=epochs,
504         validation_data=val_dataset
505     )
506
507     return policy_model, history_onl

```

```

501
502 def simulate_random_games(n):
503     game_boards = []
504     game = OthelloBoard(n)
505
506     while True:
507         moves = game.move_generator()
508         if moves == []:
509             game.pass_counter += 1
510             game.change_turn()
511
512         else:
513             game.pass_counter = 0 # reset counter
514             action = random.choice(moves)
515             board = game.make_move(action)[0]
516             game_boards.append(np.array(board))
517             game.change_turn()
518
519             if game.pass_counter == 2:
520                 reward = game.find_winner()
521                 break
522
523     return (game_boards, reward)
524
525 def save_games(n, nr_simulations):
526     all_game_boards = []
527
528     for _ in range(nr_simulations):
529         boards, actions = simulate_random_games(n)
530         all_game_boards.append(boards)
531
532     with open(f"othello_sim_boards_{nr_simulations}", "wb") as fp: #Pickling
533         pickle.dump(all_game_boards, fp)
534
535 def build_tree(n, value_model, policy_model, c, nr_simulations=1000):
536
537     game = OthelloBoard(n)
538     root = MCTSNode(n, game.board, game.to_play)
539     episodes = []
540     mapped_actions = map_actions_to_integers(n)
541
542     while True:
543         terminal_state, reward = expand_tree_iteratively(root, value_model, policy_
544 model, mapped_actions, c)
545         episodes.append((episode(terminal_state)[-1], reward))
546         print('Episode done!')
547
548         if len(episodes) == nr_simulations:
549             break
550
551     return root, episodes
552
553 def expand_tree_iteratively(root, value_model, policy_model, mapped_actions, c):
554     stack = [(root, root)] # Stack holds pairs of (root, root)
555
556     while stack:
557         current_root, current_node = stack.pop() # Get the last node to process
558
559         if current_node._untried_actions: # If the node has untried actions, expand
560 it
561             next_node = current_node.expand()
562             # Instead of using recursion, use a random value for the q_val
563             q_val = np.array(value_model(np.expand_dims(np.array(next_node.state),
564 axis=(0, -1))))[0][0])
565             next_node.update_q(q_val)

```

```

563         next_node.backpropagate(next_node.q_value)
564
565         # Push the root back into the stack to continue processing from the root
566         stack.append((current_root, current_root))
567
568     else: # If no untried actions, move to the best child or terminal node
569         if current_node.pass_counter != 2:
570
571             flatten_state = np.ndarray.flatten(np.array(current_node.state))
572             policy_dist = np.array(policy_model(np.expand_dims(flatten_state, axis
=0))[0])
573             for child in current_node.child_nodes:
574                 child.p_action = policy_dist[mapped_actions[child.parent_action]]
575
576             best_child = current_node.best_child(c)
577             stack.append((current_root, best_child)) # Continue with the best
child
578
579         else: # Terminal node/state
580             reward = current_node.find_winner(current_node.state)
581             current_node.update_q(reward)
582             current_node.backpropagate(reward)
583             current_node.terminal_visits += 1
584             return current_node, reward
585
586 def unpack_plot_history(history_list):
587     for i, history in enumerate(history_list):
588         if history != []:
589             label = f"C{i}"
590             plt.plot(history.history['loss'], label=f'Training loss', color = label)
591             plt.plot(history.history['val_loss'], linestyle="dashed", label = '
Validation loss', color = label)
592
593             plt.xlabel('Epoch')
594             plt.ylabel('Loss')
595             plt.legend(loc='upper right')
596             plt.show()
597
598 def cycle_results_histogram(episodes_list):
599     """
600     Takes in an ordered (from first to last training cycle) episodes list where the
length of the
601     list denotes the number of training cycles that were needed to create it (each
element in the
602     list should be a list of episodes).
603
604     In return, plots a histogram with sections colored according to the number of
white wins, black
605     wins, and draws (grey).
606
607     As white should win at 6x6 othello, we expect to see the bars further to the right
to be
608     increasingly dominated by the white sections.
609     """
610     results = []
611     for episodes in episodes_list:
612         results.append(results_distribution(episodes))
613
614     # Generate x-axis indices for the number of tuples in the list
615     x = range(len(results))
616     labels = list(range(1, len(results) + 1))
617
618     # Unpack each tuple into separate segments
619     segment1 = [t[0] for t in results] # Black section
620     segment2 = [t[1] for t in results] # Grey section

```

```

621 segment3 = [t[2] for t in results] # White section
622
623 plt.figure(figsize=(11, 5))
624
625 # Plot the first segment (black)
626 plt.bar(x, segment1, color='black', edgecolor='black')
627
628 # Plot the second segment (grey), stacked on top of the first
629 plt.bar(x, segment2, bottom=segment1, color='grey', edgecolor='black')
630
631 # Compute the bottom for the third segment (segment1 + segment2)
632 bottom3 = [a + b for a, b in zip(segment1, segment2)]
633
634 # Plot the third segment (white)
635 plt.bar(x, segment3, bottom=bottom3, color='white', edgecolor='black')
636
637 # # Add a text annotation above each bar showing the white percentage.
638 # for i, t in enumerate(results):
639 #     total = sum(t)
640 #     # Calculate the white percentage (t[2] is the white segment)
641 #     white_percentage = (t[2] / total) * 100 if total else 0
642 #     # Use .3g to format the number with a maximum of 3 significant digits
643 #     plt.text(i, total + 1, f'{white_percentage:.3g}%', ha='center', va='bottom',
644 #             fontsize=5)
645
646 plt.xlabel('Training cycle')
647 plt.ylabel('Episodes')
648 plt.ylim(0, max([sum(t) for t in results]) * 1.1) # Slightly higher than max for
649 # visual clarity
650 plt.xticks(x, labels)
651 plt.show()
652
653 def terminal_visits_histogram_colored(episodes):
654     """
655     Takes in a list of episodes and returns a histogram with as many bars as
656     unique terminal states that were visited in those episodes, with each bar
657     colored according to who wins (or draw) in that position. White bars for
658     white wins, black bars for black wins, and grey for draws.
659     """
660
661     nr_games = len(episodes)
662     win, draw, loss = results_distribution(episodes)
663     w_rate, d_rate, l_rate = np.round(np.array([win, draw, loss])*(100/nr_games),
664     decimals=0)
665
666     terminal_states = {}
667     for episode in episodes:
668         key = (episode[0][0].tobytes(), episode[1])
669         if key not in terminal_states:
670             terminal_states[key] = 1
671         else:
672             terminal_states[key] += 1
673
674     color_mapping = {
675         -1: 'white', # white for -1
676         0: 'grey', # grey for 0
677         1: 'black' # black for 1
678     }
679
680     # Extract rewards, visits, and colors
681     identifiers = []
682     heights = []
683     colors = []
684     for (identifier, y_value), height in terminal_states.items():
685         if isinstance(identifier, bytes):

```



```

683         identifier = identifier.decode('latin-1')
684         identifiers.append(identifier)
685         heights.append(height)
686         colors.append(color_mapping[y_value])
687
688     x = range(len(terminal_states))
689
690     plt.figure(figsize=(7.2, 5))
691     plt.bar(identifiers, heights, color=colors, edgecolor='black')
692
693     # Add a text box with game results
694     textstr = f"Total games = {nr_games}\nBlack wins: {w_rate}%\nDraws: {d_rate}%\n\nWhite wins: {l_rate}%"
695
696     # Position the text box
697     props = dict(boxstyle='round', facecolor='white', alpha=0.8)
698     plt.text(0.025, 0.975, textstr, transform=plt.gca().transAxes, fontsize=9,
699             verticalalignment='top', bbox=props)
700
701
702     plt.xticks(x, [str(i + 1) for i in x])
703
704     plt.gca().yaxis.set_major_locator(mticker.MaxNLocator(integer=True))
705
706     plt.xlabel('Unique terminal states')
707     plt.ylabel('Visits')
708
709     plt.show()
710
711 def remove_empty_lists(list):
712     # remove all empty lists from the list
713     list = [x for x in list if x]
714     return list
715
716 def unpack_plot_train_val_curves(history_list, cycle_labels):
717     """
718     All empty lists have to be removed from the input list here in order for it to work.
719     See above helper function.
720
721     Cycle_labels should be a list of labels for the cycles, i.e. [3,6,9,...] if
722     every third cycle is passed in the history list
723     """
724
725     fig = plt.figure(figsize=(12, 6))
726
727     for i, history in enumerate(history_list):
728         plt.subplot(2,5,i+1)
729         color = f"C{i}"
730
731         plt.plot(history.history['loss'], label=f'Training loss', color = color)
732         plt.plot(history.history['val_loss'], linestyle="dashed", label = 'Validation
733         loss', color = color)
734         plt.title(f'Cycle {cycle_labels[i]}')
735         plt.tick_params(axis='x', which='both', bottom=False,
736             top=False, labelbottom=False)
737         plt.tick_params(axis='y', which='both', right=False,
738             left=False, labelleft=False)
739
740         textstr = '\u2500 * 5 + ' Training loss \n----- Validation loss'
741         props = dict(boxstyle='round', facecolor='white', alpha=0.8)
742         plt.text(-0.125, -0.10, textstr, transform=plt.gca().transAxes, fontsize=12,
743             verticalalignment='top', bbox=props)
744
745     fig.text(0.5, 0.04, 'Epoch', va='center', ha='center', fontsize=14)

```

```

745     fig.text(0.09, 0.5, 'Loss', va='center', ha='center', rotation='vertical',
746             fontsize=14)
747     plt.show()
748     #####
749     # ----- Beginning of VALUE NN architecture -----
750     #####
751
752     def residual_block(x, channels=64, kernel_size=(3,3), weight_decay=0.001):
753         shortcut = x # No projection needed if dimensions already match.
754
755         x = layers.Conv2D(channels, kernel_size=kernel_size,
756                           padding='same', use_bias=False,
757                           kernel_regularizer=regularizers.l2(weight_decay))(x)
758         x = layers.BatchNormalization()(x)
759         x = layers.ReLU()(x)
760
761         x = layers.Conv2D(channels, kernel_size=kernel_size,
762                           padding='same', use_bias=False,
763                           kernel_regularizer=regularizers.l2(weight_decay))(x)
764         x = layers.BatchNormalization()(x)
765
766         # Direct addition is fine here.
767         x = layers.Add()([x, shortcut])
768         x = layers.ReLU()(x)
769
770         return x
771
772     # Build the model
773     inputs = layers.Input(shape=(6, 6, 1)) # 6x6 Othello
774
775     # Initial convolution block (producing 16 channels)
776     x = layers.Conv2D(64, kernel_size=(3,3),
777                       padding='same', use_bias=False,
778                       kernel_regularizer=regularizers.l2(0.001))(inputs)
779     x = layers.BatchNormalization()(x)
780     x = layers.ReLU()(x)
781
782     # Add a number of residual blocks with constant 64 channels
783     x = residual_block(x, channels=64, kernel_size=(3,3), weight_decay=0.001)
784     x = residual_block(x, channels=64, kernel_size=(3,3), weight_decay=0.001)
785
786     # Flatten the features and output a single scalar value (for a value network)
787     x = layers.Flatten()(x)
788     outputs = layers.Dense(1, name="value_output")(x)
789
790     # Create the model
791     value_model = Model(inputs=inputs, outputs=outputs)
792     value_model.summary()
793
794     # Optimizer for the value model
795     value_model.compile(
796         optimizer=keras.optimizers.SGD(learning_rate=1e-4, momentum=0.9),
797         loss = keras.losses.MeanSquaredError()
798         # metrics = [keras.metrics.RootMeanSquaredError]
799     )
800
801     #####
802     # ----- Beginning of Policy NN architecture -----
803     #####
804
805     # Policy model architecture
806     policy_model = keras.Sequential([
807
808         layers.Dense(128, input_shape=(36,)),

```

```

809     layers.BatchNormalization(),
810     layers.Activation("relu"),
811     layers.Dropout(0.025),
812
813     layers.Dense(256),
814     layers.BatchNormalization(),
815     layers.Activation("relu"),
816     layers.Dropout(0.025),
817
818     layers.Dense(256),
819     layers.BatchNormalization(),
820     layers.Activation('relu'),
821     layers.Dropout(0.025),
822
823     layers.Dense(256),
824     layers.BatchNormalization(),
825     layers.Activation("relu"),
826     layers.Dropout(0.025),
827
828     layers.Dense(128),
829     layers.BatchNormalization(),
830     layers.Activation('relu'),
831     layers.Dropout(0.025),
832
833     layers.Dense(36),
834     layers.Activation('softmax')
835 ])
836
837 policy_model.summary()
838
839 policy_model.compile(optimizer=keras.optimizers.SGD(learning_rate=1e-3, momentum=0.9),
840                     loss=keras.losses.KLDivergence())
841
842 # %%
843 # Loading the model weights
844 # Load the model weights
845 # value_model.load_weights('zeroth_value_nn.weights.h5')
846 # policy_model.load_weights('zeroth_policy_nn.weights.h5')
847
848
849 #####
850 # ----- Online training section -----
851 #####
852
853 """
854 Given a replay buffer, we want to be able to continuously feed new game
855 information into the value network in the form of mini batches. This
856 section aims to prepare for that.
857
858 Replay buffer will be in the form of a list of tuples, where the first
859 elements in each tuple is a game consisting of a sequence of board positions
860 and the second element is the observed reward from that game.
861
862 Prerequisites for the following function is to already have pre-trained
863 value and policy networks.
864 """
865
866 # %%
867 def online_simulation(n, c, value_model, policy_model, nr_cycles, nr_episodes_per_tree
868                     = 160, history_val = [], history_pol = []):
869     ''' hehiha '''
870
871     trees = []
872     episodes = []

```

```

873 value_model_history = [history_val]
874 policy_model_history = [history_pol]
875
876 for _ in range(nr_cycles):
877
878     ##### Build tree with updated model #####
879     tree, replay_buffer = build_tree(n, value_model, policy_model, c, nr_episodes_
per_tree)
880
881     trees.append(tree)
882     episodes.append(replay_buffer)
883
884     ##### Update value network #####
885     value_model, value_history_temp = online_value_training(value_model, replay_
buffer,
886
887
888                                     epochs=20, batch_size
=64)
889     policy_model, policy_history_temp = online_policy_training(policy_model, tree,
890
891                                     epochs=50, batch_size
=128)
892     value_model_history.append(value_history_temp)
893     policy_model_history.append(policy_history_temp)
894
895     print('Cycle done!')
896
897
898 return trees, episodes, value_model_history, policy_model_history
899
900 # %%
901 trees, episodes, val_hist, pol_hist = online_simulation(n=6, c=0.5,
902
903                                     value_model=value_model,
904                                     policy_model=policy_model,
905                                     nr_cycles=10,
906                                     nr_episodes_per_tree=25)
907
908 # %%
909 # saving the training from the online training
910 second_10cycles25episodes_c05_data = [trees, episodes, val_hist, pol_hist]
911 with open('second_10cycles25episodes_c05_data', 'wb') as handle:
912     pickle.dump(second_10cycles25episodes_c05_data,
913                 handle, protocol=pickle.HIGHEST_PROTOCOL)
914
915 value_model.save_weights('value_second_10cycles25episodes_c05.weights.h5')
916 policy_model.save_weights('policy_second_10cycles25episodes_c05.weights.h5')
917
918 #####
919 # ----- Initial training of VALUE NN -----
920 #####
921
922 # Initial value network data - from completely random games
923 # Unique board positions from first 100k random games
924 path = './othello_random_simulations/othello_sim_boards_100000_6x6'
925 games = read_files(path)
926
927 # %%
928 state_return_tuples = unpack_positions_returns(games)
929 state_return_tuples = unique_board_positions(state_return_tuples)
930
931 X, y = value_predictors_targets(state_return_tuples)
932
933 # splitting data into training and validation
934 X_train_val, X_test_val, y_train_val, y_test_val = train_test_split(X, y, test_size
=0.2,

```

```

932                                     random_state
                                     =80085)
933
934 train_dataset_value_model = create_dataset(X_train_val, y_train_val, batch_size=256,
935                                           shuffle_buffer_size=10000)
936 val_dataset_value_model = create_dataset(X_test_val, y_test_val, batch_size=256,
937                                         shuffle_buffer_size=10000)
938
939 # Training the value model
940 history_value = value_model.fit(
941     train_dataset_value_model,
942     epochs = 50,
943     validation_data=val_dataset_value_model
944     # callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)]
945 )
946
947 # with open('initial_value_training_history', 'wb') as handle:
948 #     pickle.dump(history_value, handle, protocol=pickle.HIGHEST_PROTOCOL)
949
950 # # Saving the weights from the first good network
951 # value_model.save_weights('zeroth_value_nn.weights.h5')
952
953 # %%
954 #####
955 # ----- Initial training of Policy NN -----
956 #####
957
958 sav = read_files("./othello_random_simulations/othello_sim_sa_visits")
959 X, y = policy_predictors_targets(sav, 6)
960
961 X_train_pol, X_test_pol, y_train_pol, y_test_pol = train_test_split(X, y, test_size
962                             =0.2,
963                             random_state=80085)
964
965 train_dataset_policy_model = create_dataset(X_train_pol, y_train_pol, batch_size=128)
966 val_dataset_policy_model = create_dataset(X_test_pol, y_test_pol, batch_size=128)
967
968 history_policy = policy_model.fit(train_dataset_policy_model,
969                                  validation_data=val_dataset_policy_model,
970                                  epochs=500)
971
972 # with open('initial_policy_training_history_2', 'wb') as handle:
973 #     pickle.dump(history_policy, handle, protocol=pickle.HIGHEST_PROTOCOL)
974
975 # saving weights from first good network
976 # policy_model.save_weights('zeroth_policy_nn_2.weights.h5')
977
978 #####
979 # ----- Concatenating history from different runs -----
980 #####
981
982 # concatenating and plotting history from first and second iteration of VALUE network
983 initial_value_training_history_1 = read_files('initial_value_training_history')
984 initial_value_training_history_2 = read_files('initial_value_training_history_2')
985
986 complete_value_history_train_loss = initial_value_training_history_1.history['loss'] +
987     initial_value_training_history_2.history['loss']
988
989 complete_value_history_val_loss = initial_value_training_history_1.history['val_loss']
990     + initial_value_training_history_2.history['val_loss']
991
992 plt.plot(complete_value_history_train_loss, label='Training loss')
993 plt.plot(complete_value_history_val_loss, label = 'Validation loss')
994 plt.xlabel('Epoch')
995 plt.ylabel('Loss')

```

```

993 plt.title('value network')
994 plt.legend(loc='upper right')
995 plt.show()
996
997 # %%
998 # concatenating and plotting history from first and second iteration of POLICY network
999 initial_policy_training_history_1 = read_files('initial_policy_training_history')
1000 initial_policy_training_history_2 = read_files('initial_policy_training_history_2')
1001
1002 complete_policy_history_train_loss = initial_policy_training_history_1.history['loss']
      + initial_policy_training_history_2.history['loss']
1003 complete_policy_history_val_loss = initial_policy_training_history_1.history['val_loss']
      + initial_policy_training_history_2.history['val_loss']
1004
1005 plt.plot(complete_policy_history_train_loss, label='Training loss')
1006 plt.plot(complete_policy_history_val_loss, label = 'Validation loss')
1007 plt.xlabel('Epoch')
1008 plt.ylabel('Loss')
1009 plt.title('policy network')
1010 plt.legend(loc='upper right')
1011 plt.show()

```